

Informed Search Methods

Sunil Surve

Fr. Conceicao Rodrigues College of Engineering

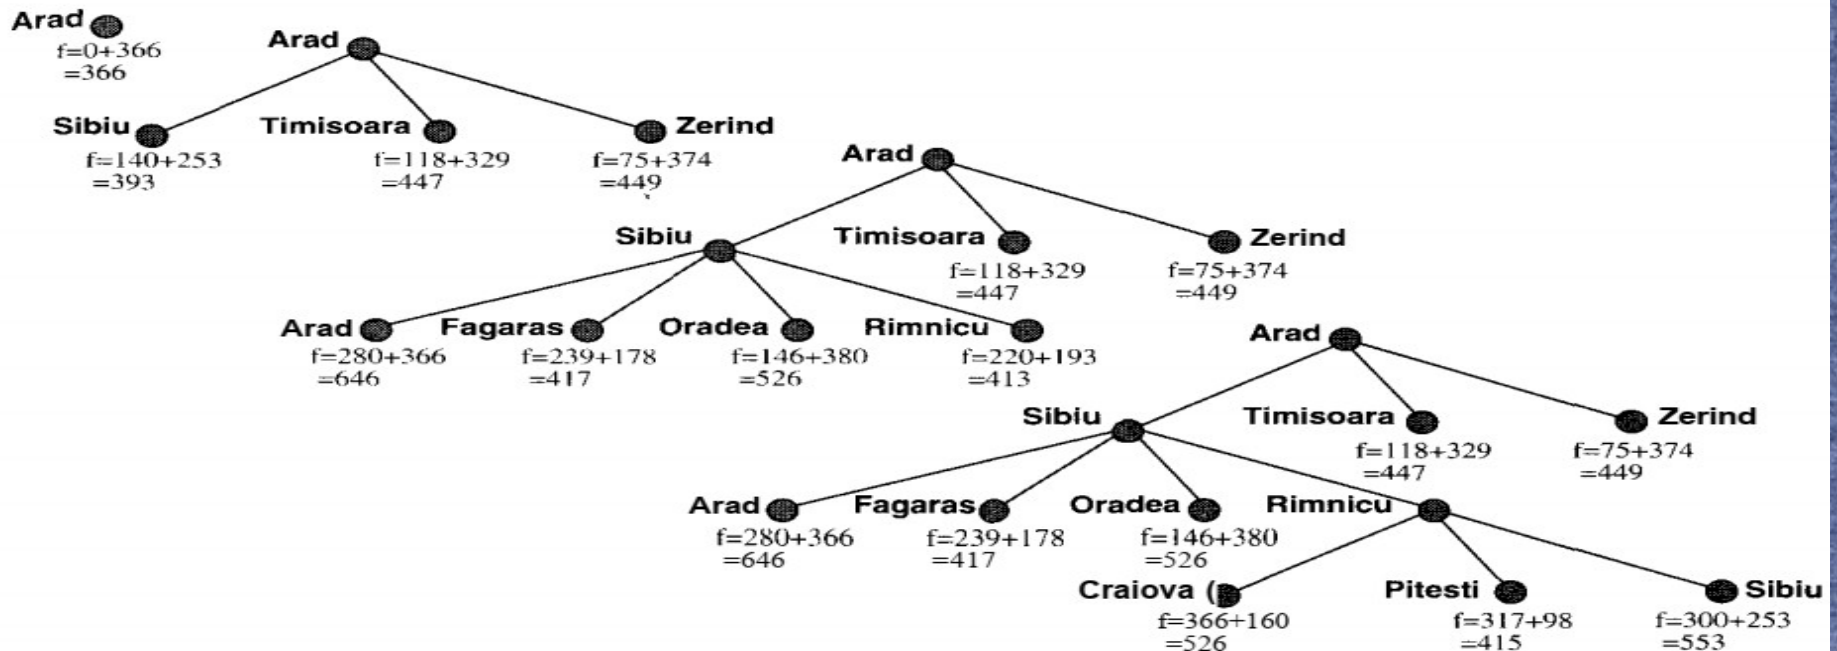
A* Algorithm

- Greedy search minimizes the estimated cost to the goal, $h(n)$, and thereby cuts the search cost considerably
- Uniform-cost search, on the other hand, minimizes the cost of the path so far, $g(n)$; it is optimal and complete, but can be very inefficient
- Combining the two evaluation functions simply by summing them:
$$f(n) = g(n) + h(n)$$

where $g(n)$ is the path cost from the start node to node n , $h(n)$ is the estimated cost of the cheapest path from n to the goal and $f(n)$ is estimated cost of the cheapest solution through n
- **Admissible heuristic**
- *If h is admissible, $f(n)$ never overestimates the actual cost of the best solution through n*


A* Algorithm

- Along any path from the root, the f-cost never decreases.
- It holds true for almost all admissible heuristics.
- A heuristic for which it holds is said to exhibit **monotonicity**



A* Algorithm

- If f-cost of child is less than parent's f-cost, then heuristic used is nonmonotonic.
- Pathmax equation $f(n') = \max(f(n), g(n') + h(n'))$
- If/ never decreases along any path out from the root, we can conceptually draw contours in the state space




A* Algorithm

- If we define f^* to be the cost of the optimal solution path, then we can say the following:
 - A* expands all nodes with $f(n) < f^*$.
 - A* may then expand some of the nodes right on the "goal contour," for which $f(n) = f^*$, before selecting a goal node
- First solution found must be the optimal one, because nodes in all subsequent contours will have higher/-cost, and thus higher g-cost
- A* is **optimally efficient** for any given heuristic function

```
function A*-SEARCH( problem) returns a solution or failure  
  return BEST-FIRST-SEARCH(problem,  $g + h$ )
```


Heuristic Functions

- 8-puzzle - the object of the puzzle is to slide the tiles horizontally or vertically into the empty space until the initial configuration matches the goal configuration
- The branching factor is about 3, and $3^{20} = 3.5 \times 10^9$ states and By keeping track of repeated states, there are only $9! = 362,880$ states



Heuristic Functions

- h_1 = the number of tiles that are in the wrong position. For Figure 4.7, none of the 8 tiles is in the goal position, so the start state would have $h_1 = 8$. h_1 is an admissible heuristic
- h_2 = the sum of the distances of the tiles from their goal positions. This is sometimes called the **city block distance** or **Manhattan distance**. h_2 is also admissible $h_2 = 2 + 3 + 2 + 1 + 2 + 2 + 1 + 2 = 15$
- Which heuristic function is better? h_1 or h_2 ?
- h_2 will expand lesser nodes than h_1 .

Inventing heuristic functions

- Relaxed problem
- If a collection of admissible heuristics h_1, \dots, h_m is available for a problem, and none of them dominates any of the others, which should we choose?
$$h(n) = \max(h_1(n), \dots, h_m(n)).$$
- *Another way to invent a good heuristic is to use statistical information*

Iterative deepening A* search (IDA*)

- Use an f-cost limit
- Each iteration expands all nodes inside the contour for the current f-cost
- Once the search inside a given contour has been completed, a new iteration is started using a new f-cost for the next contour
- IDA* is complete and optimal with the same caveats as A* search IDA* is complete and optimal with the same caveats as A* search
- Requires space proportional to the longest path that it explores.
- If b is the smallest operator cost and f^* the optimal solution cost, then in the worst case, IDA* will require $bf^*/8$ nodes of storage.
- In most cases, bd is a good estimate of the storage requirements

Iterative deepening A* search (IDA*)

- IDA* has difficulty in more complex domains
- In the travelling salesperson problem, for example, the heuristic value is different for every state. This means that each contour only includes one more state than the previous contour.
- If A* expands N nodes, IDA* will have to go through N iterations and will therefore expand $1 + 2 + \dots + N = O(N^2)$ nodes.
- One way around this is to increase the f-cost limit by a fixed amount ϵ on each iteration, so that the total number of iterations is proportional to $1/\epsilon$.
- This can reduce the search cost, at the expense of returning solutions that can be worse than optimal by at most ϵ . Such an algorithm is called **ϵ -admissible**

Iterative deepening A* search (IDA*)

function IDA*(*problem*) **returns** a solution sequence

inputs: *problem*, a problem

static: *f-limit*, the current *f*- COST limit

mot, a node

root \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

f-limit \leftarrow *f*- COST(*root*)

loop do

solution, *f-limit* \leftarrow DFS-CONTOUR(*root*, *f-limit*)

if *solution* is non-null **then return** *solution*

if *f-limit* = ∞ **then return** failure; **end**

function DFS-CONTOUR(*node*, *f-limit*) **returns** a solution sequence and a new *f*- COST limit

inputs: *node*, a node

f-limit, the current *f*- COST limit

static: *next-f*, the *f*- COST limit for the next contour, initially ∞

if *f*- COST[*node*] > *f-limit* **then return** null, *f*- COST[*node*]

if GOAL-TEST[*problem*](STATE[*node*]) **then return** *node*, *f-limit*

for each node *s* in SUCCESSORS(*node*) **do**

solution, *new-f* \leftarrow DFS-CONTOUR(*s*, *f-limit*)



if *solution* is non-null **then return** *solution*, *f-limit*

next-f \leftarrow MIN(*next-f*, *new-f*); **end**

return null, *next-f*

Simplified Memory-Bounded A*

- SMA* has the following properties:
 - It will utilize whatever memory is made available to it.
 - It avoids repeated states as far as its memory allows.
 - It is complete if the available memory is sufficient to store the shallowest solution path.
 - It is optimal if enough memory is available to store the shallowest optimal solution path.
 - Otherwise, it returns the best solution that can be reached with the available memory.
 - When enough memory is available for the entire search tree, the search is optimally efficient



Simplified Memory-Bounded A*

```
function SMA*(problem) returns a solution sequence
  inputs: problem, a problem
  static: Queue, a queue of nodes ordered by f-cost

  Queue ← MAKE-QUEUE([ MAKE-NODE(INITIAL-STATE[problem])) )
  loop do
    if Queue is empty then return failure
    n ← deepest least-f-cost node in Queue
    if GOAL-TEST(n) then return success
    s ← NEXT-SUCCESSOR(n)
    if s is not a goal and is at maximum depth then
       $f(s) \leftarrow \infty$ 
    else
       $f(s) \leftarrow \text{MAX}(f(n), g(s)+h(s))$ 
    if all of n's successors have been generated then
      update n's f-cost and those of its ancestors if necessary
    if SUCCESSORS(n) all in memory then remove n from Queue
    if memory is full then
      delete shallowest, highest-f-cost node in Queue
      remove it from its parent's successor list
      insert its parent on Queue if necessary
    insert s on Queue
  end
```


Hill-climbing search

- A loop that continually moves in the direction of increasing value.
- The algorithm does not maintain a search tree
- Local maxima: a local maximum is a peak that is lower than the highest peak in the state space. Once on a local maximum, the algorithm will halt even though the solution may be far from satisfactory
- Plateaux: a plateau is an area of the state space where the evaluation function is essentially flat.
- Ridges: a ridge may have steeply sloping sides, so that the search reaches the top of the ridge with ease, but the top may slope only very gently toward a peak

