

# Map-Reduce

## Learning Objectives:

Distributed File System Basics.

Challenges


Why Map-Reduce?

## Mining of Massive Datasets

J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets, <http://www.mmds.org>




# Single Node Architecture



**Machine Learning, Statistics**

**“Classical” Data Mining**

# Single Node Architecture




**Machine Learning, Statistics**

**“Classical” Data Mining**


# Motivation: Google Example

- ‡ 10 billion web pages
- ‡ Average size of webpage = 20KB
- ‡ 10 billion \* 20KB = 200 TB
- ‡ Disk read bandwidth = 50 MB/sec
- ‡ Time to read = 4 million seconds = 46+ days
- ‡ Even longer to do something useful with the data



J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets, <http://www.mmids.org>

# Cluster Architecture



Each rack contains 16-64 commodity Linux nodes

In 2011 it was guestimated that Google had 1M machines, <http://bit.ly/Shh0RO>

J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets, <http://www.mmms.org>



# Cluster Computing Challenges (1)

## ➤ Node failures

§ A single server can stay up for 3 years (1000 days)

§ 1000 servers in cluster => 1 failure/day

§ 1M servers in cluster => 1000 failures/day

➤ How to store data **persistently** and keep it **available** if nodes can fail?

➤ How to deal with node failures during a long-running computation?

# Cluster Computing Challenges (2)

## ➤ Network bottleneck

§ Network bandwidth = 1 Gbps

§ Moving 10TB takes approximately 1 day

## ➤ Distributed programming is hard!

§ Need a simple model that hides most of the complexity



# Map-Reduce

- Map-Reduce addresses the challenges of cluster computing
  - § Store data redundantly on multiple nodes for persistence and availability
  - § Move computation close to data to minimize data movement
  - § Simple programming model to hide the complexity of all this distributed computing.

# Redundant Storage Infrastructure

## Distributed File System

§ Provides global file namespace, redundancy, and availability

§ E.g., Google GFS; Hadoop HDFS

## Typical usage pattern

§ Huge files (100s of GB to TB)

§ Data is rarely updated in place


§ Reads and appends are common

# Distributed File System

Data kept in “chunks” spread across machines

Each chunk **replicated** on different machines

§ Ensures persistence and availability



Chunk servers also serve as compute servers

Bring computation to data!

# Distributed File System

## Chunk servers

- § File is split into contiguous chunks (16-64MB)
- § Each chunk replicated (usually 2x or 3x)
- § Try to keep replicas in different racks

## Master node

- § a.k.a. Name Node in Hadoop's HDFS
- § Stores metadata about where files are stored
- § Might be replicated

## Client library for file access

- § Talks to master to find chunk servers
- § Connects directly to chunk servers to access data